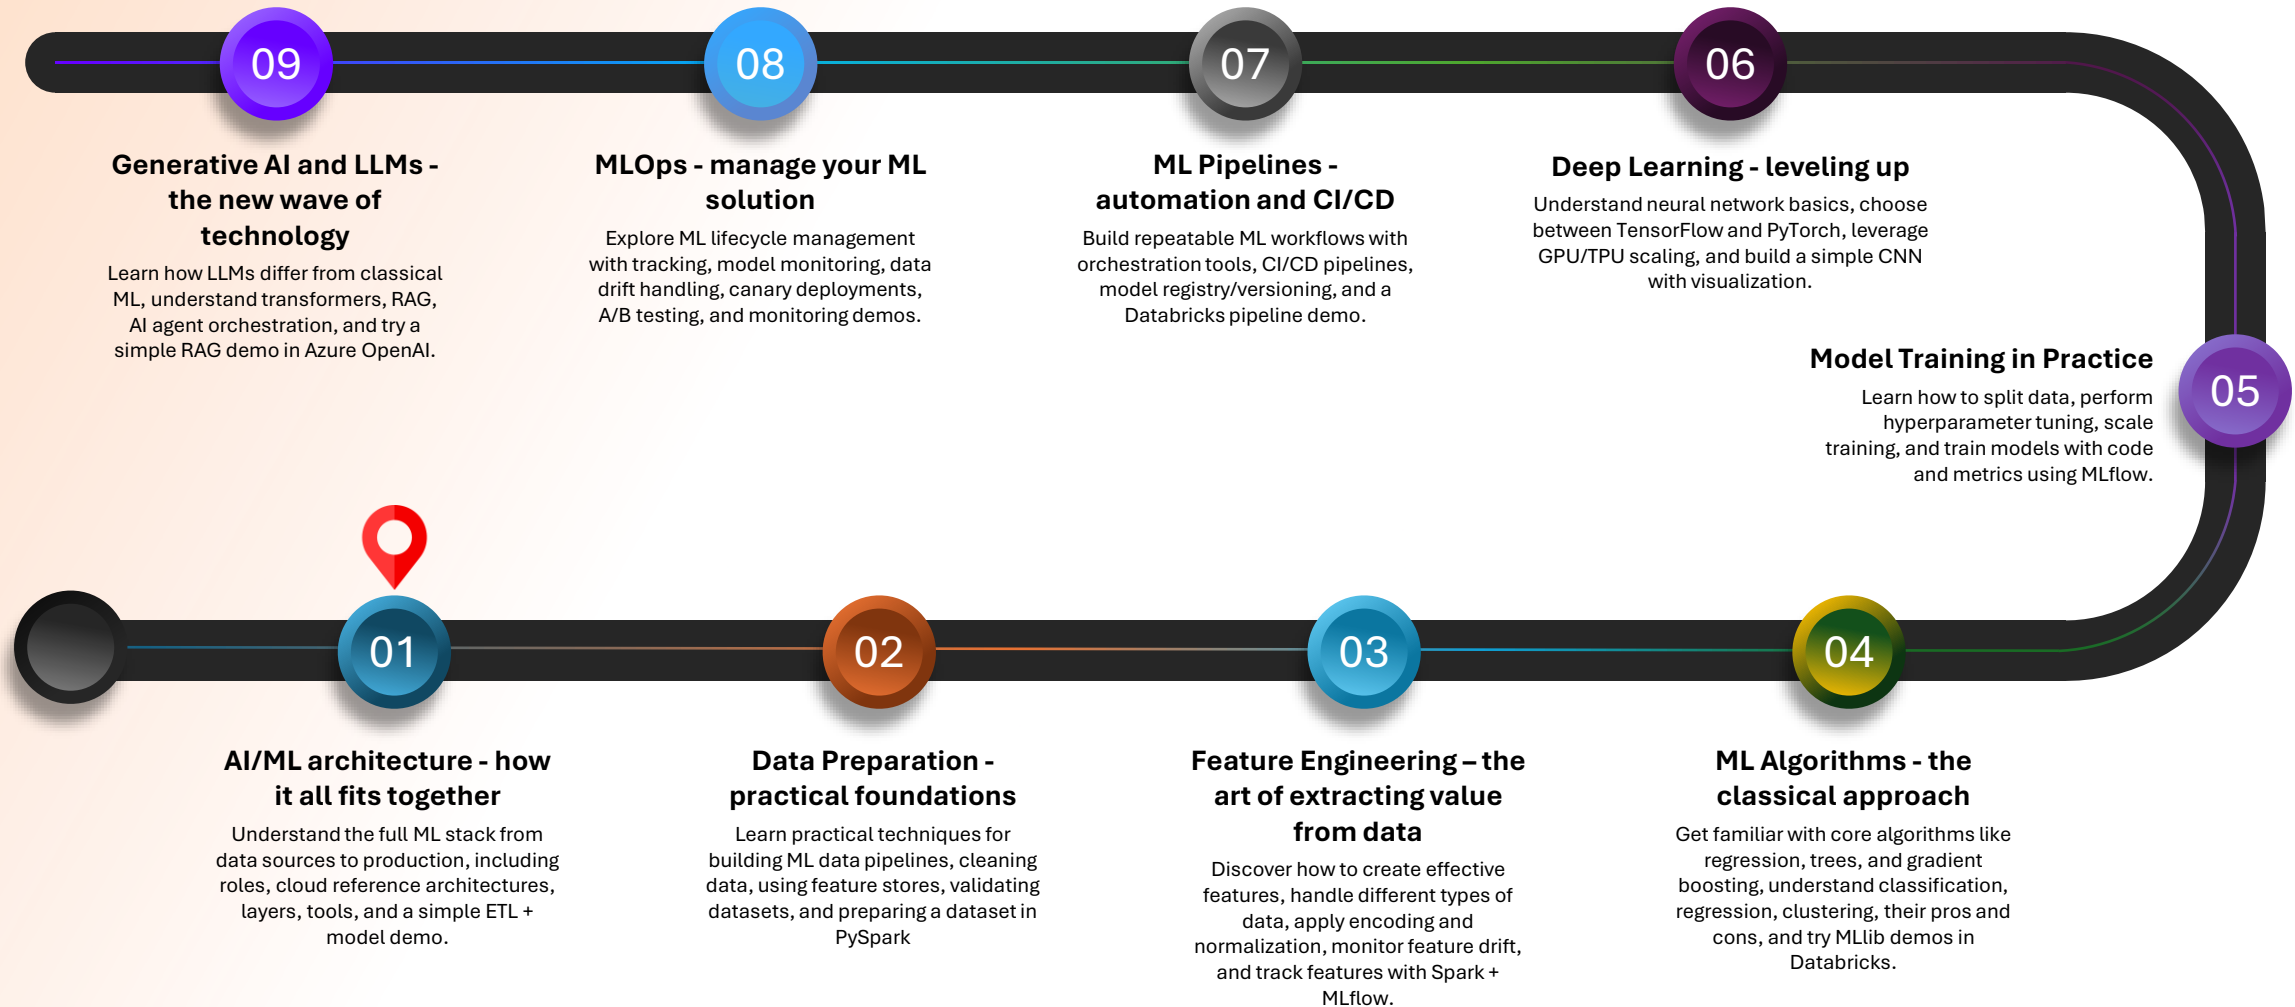


# **AI/ML Architecture - how it all fits together**

**Maciej Kępa**

# Roadmap



# Agenda

1. Introduction
2. What an AI/ML system is made of?
3. Roles and responsibilities
4. Cloud architecture
5. Demo
6. System design best practices


# AI history



Long, long ago...




1950s



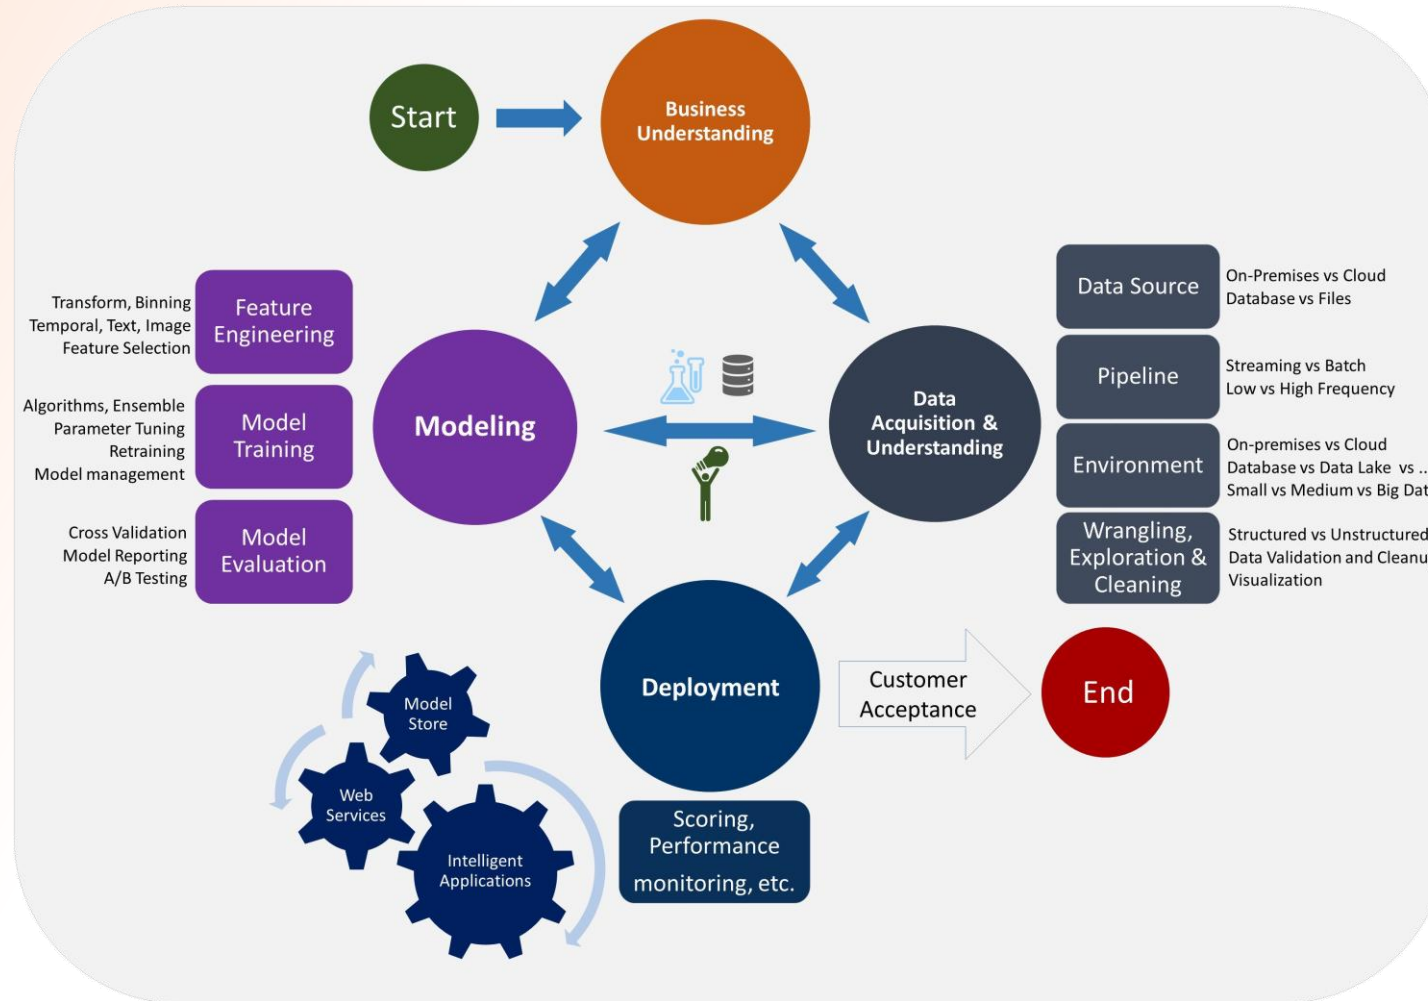
2017

Completely **A**utomated **P**ublic **T**uring test to tell **C**omputers and **H**umans **A**part.

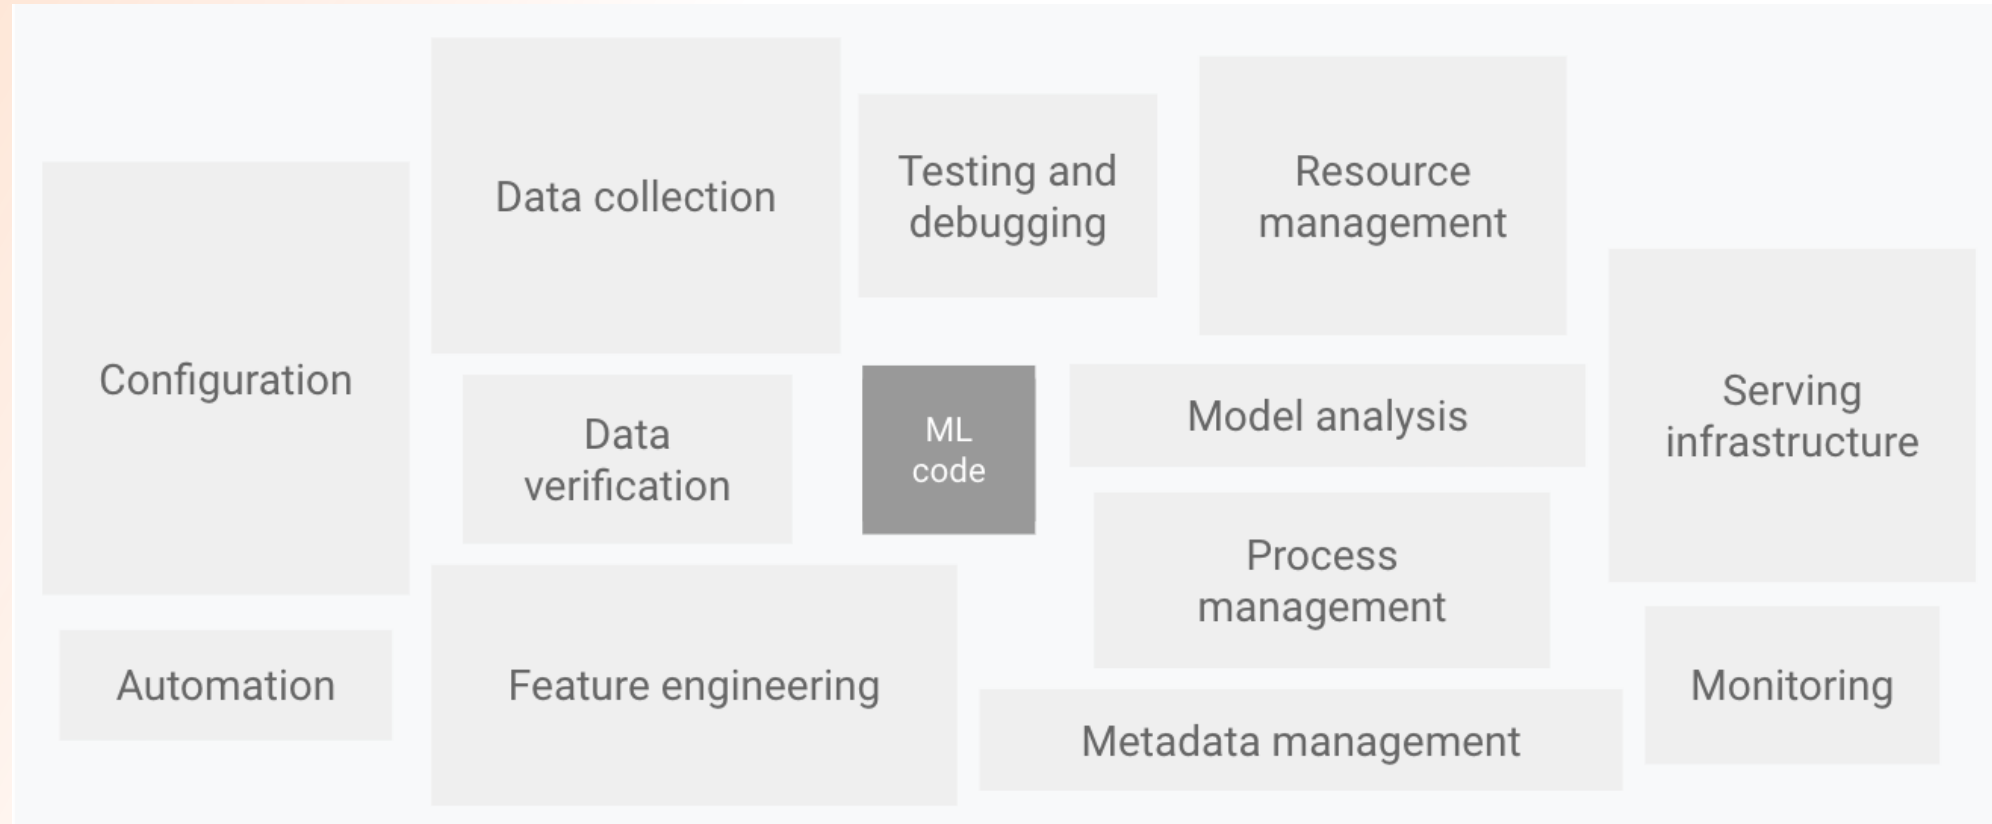


**What an AI/ML  
system is made of?**

# AI project lifecycle




# AI project areas



# Key layers

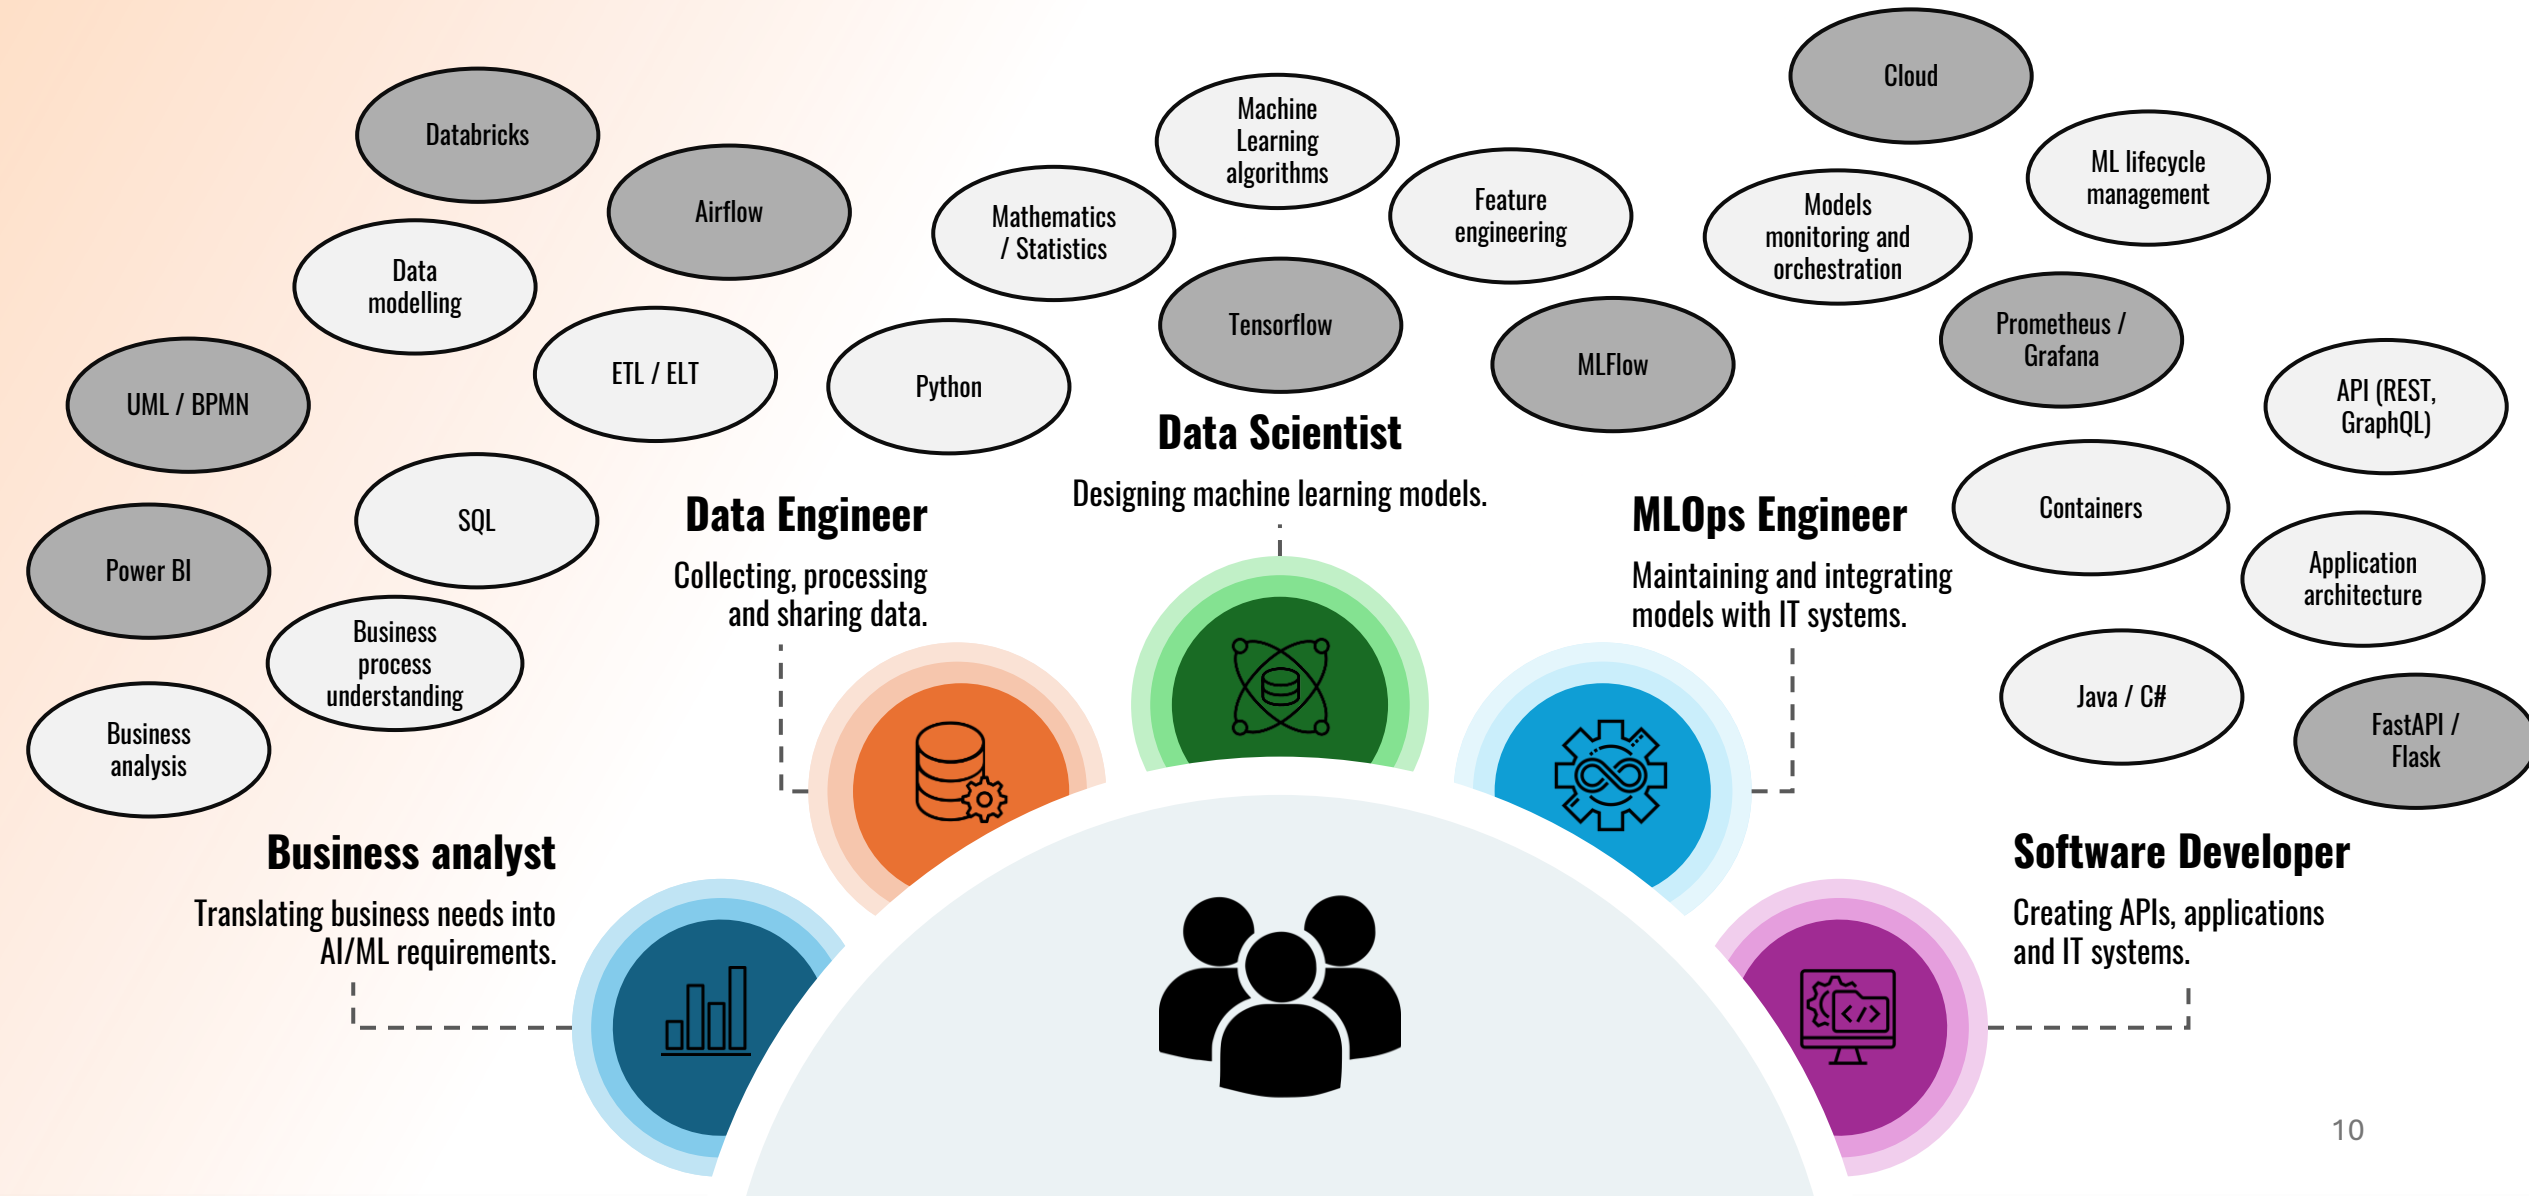
Layer	Purpose	Typical Components / Tools	Key Responsibilities	Artifacts
<b>Data Layer</b>	Collects, stores, and prepares data for modeling.	Data Lake (Azure Data Lake, GCS, S3), Data Warehouse (Synapse, BigQuery), Feature Store (Databricks Feature Store, Feast).	Ingestion from raw sources, cleaning, transformation, feature creation, and validation (Great Expectations).	Raw and processed datasets, feature tables, data validation reports, ETL notebooks, schema definitions.
<b>Model Layer</b>	Focuses on training, tracking, and versioning ML models.	Databricks, MLflow, Azure ML, TensorFlow, PyTorch, scikit-learn.	Model development, experiment tracking, hyperparameter tuning, model registration, and reproducibility.	Model training notebooks, experiment logs, metrics, model binaries, registered models.
<b>Deployment Layer</b>	Delivers trained models to production and maintains operational quality.	CI/CD pipelines (GitHub Actions, Azure DevOps), Model Registry, APIs, Monitoring (Azure Monitor, Prometheus).	Model packaging and deployment, version control, A/B testing, drift monitoring, continuous retraining.	Deployment scripts, inference APIs, dashboards, monitoring reports, retraining jobs.






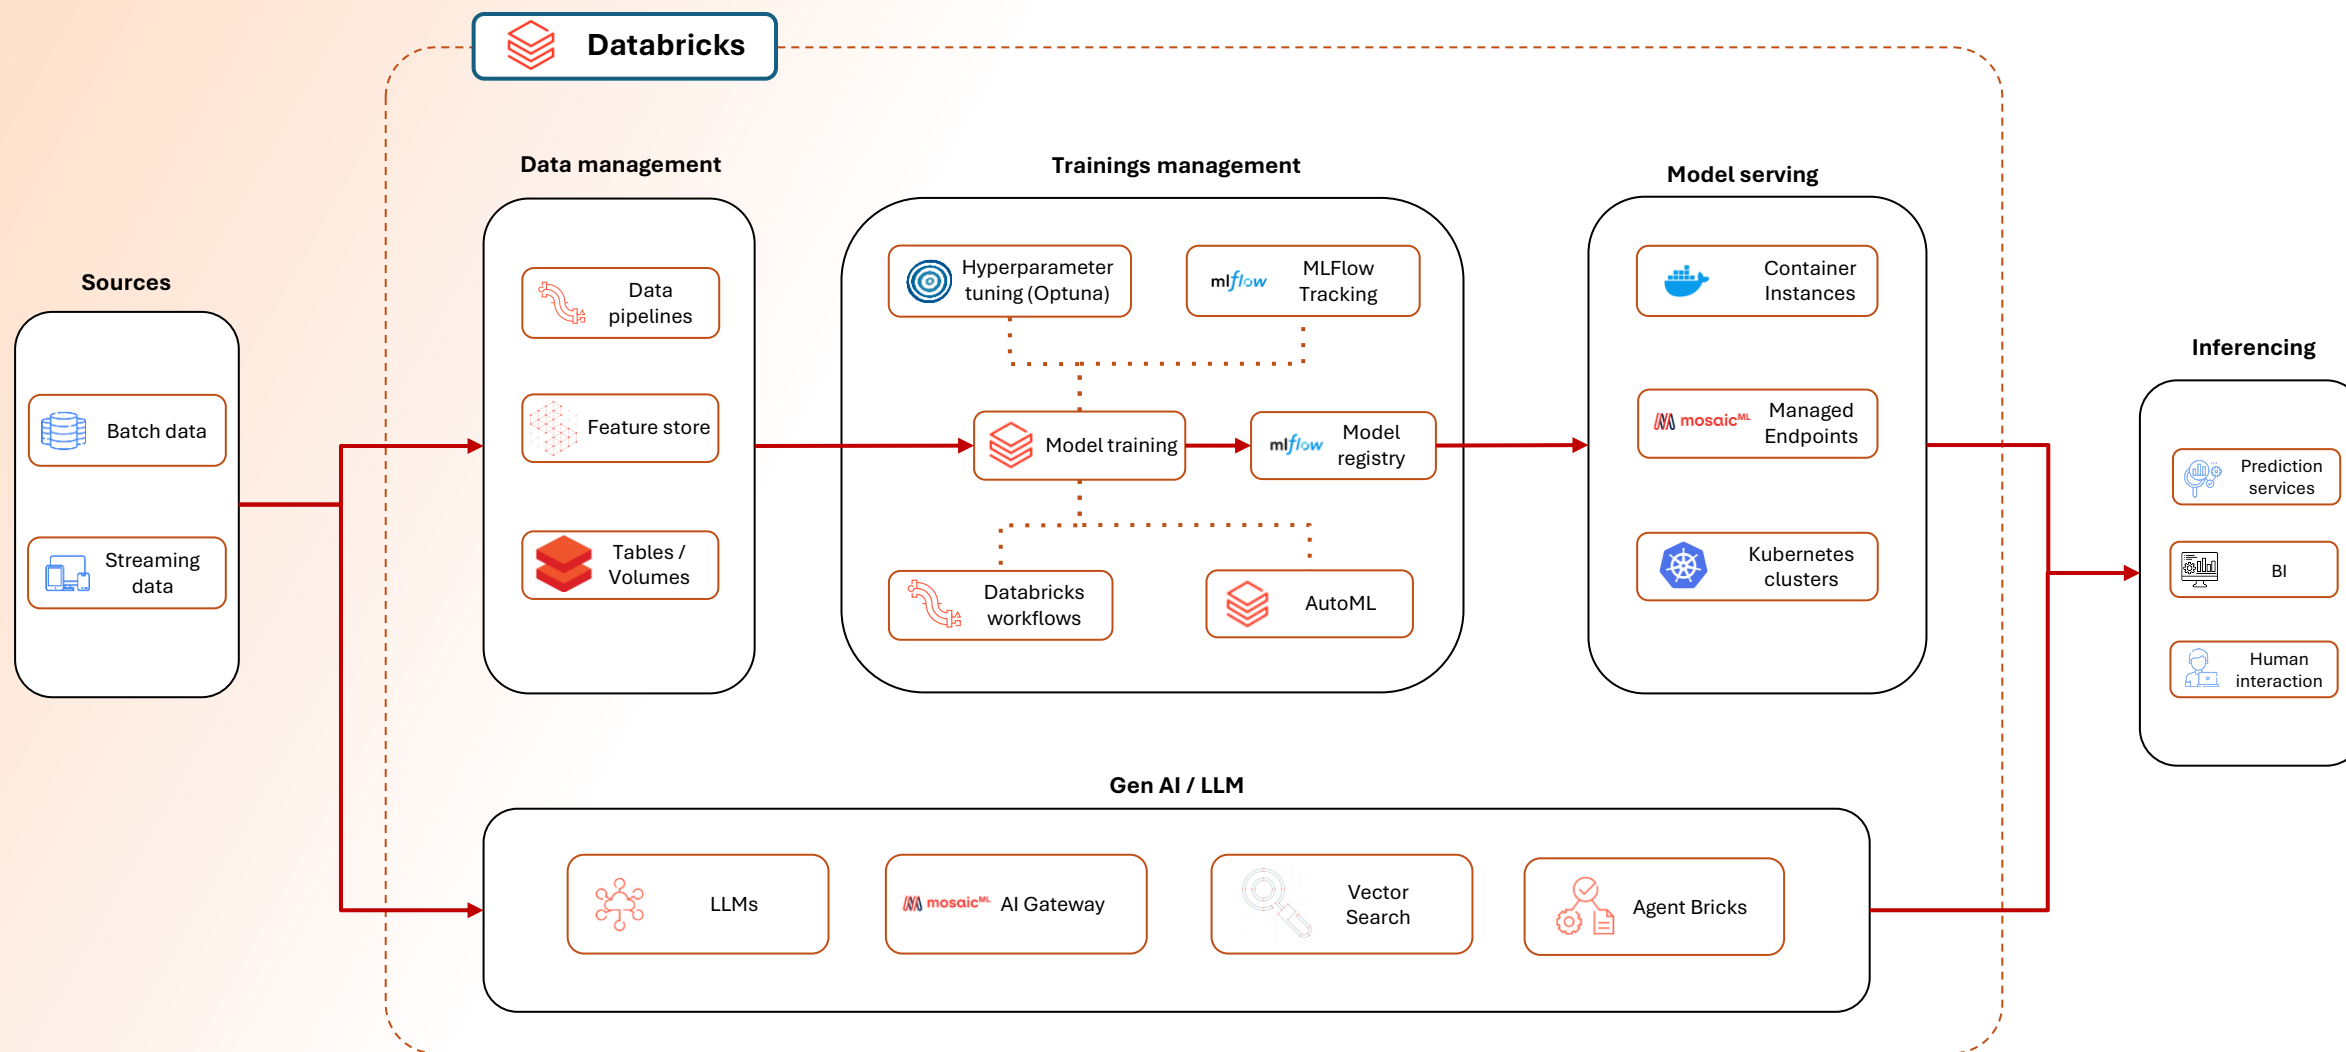
# Roles and responsibilities

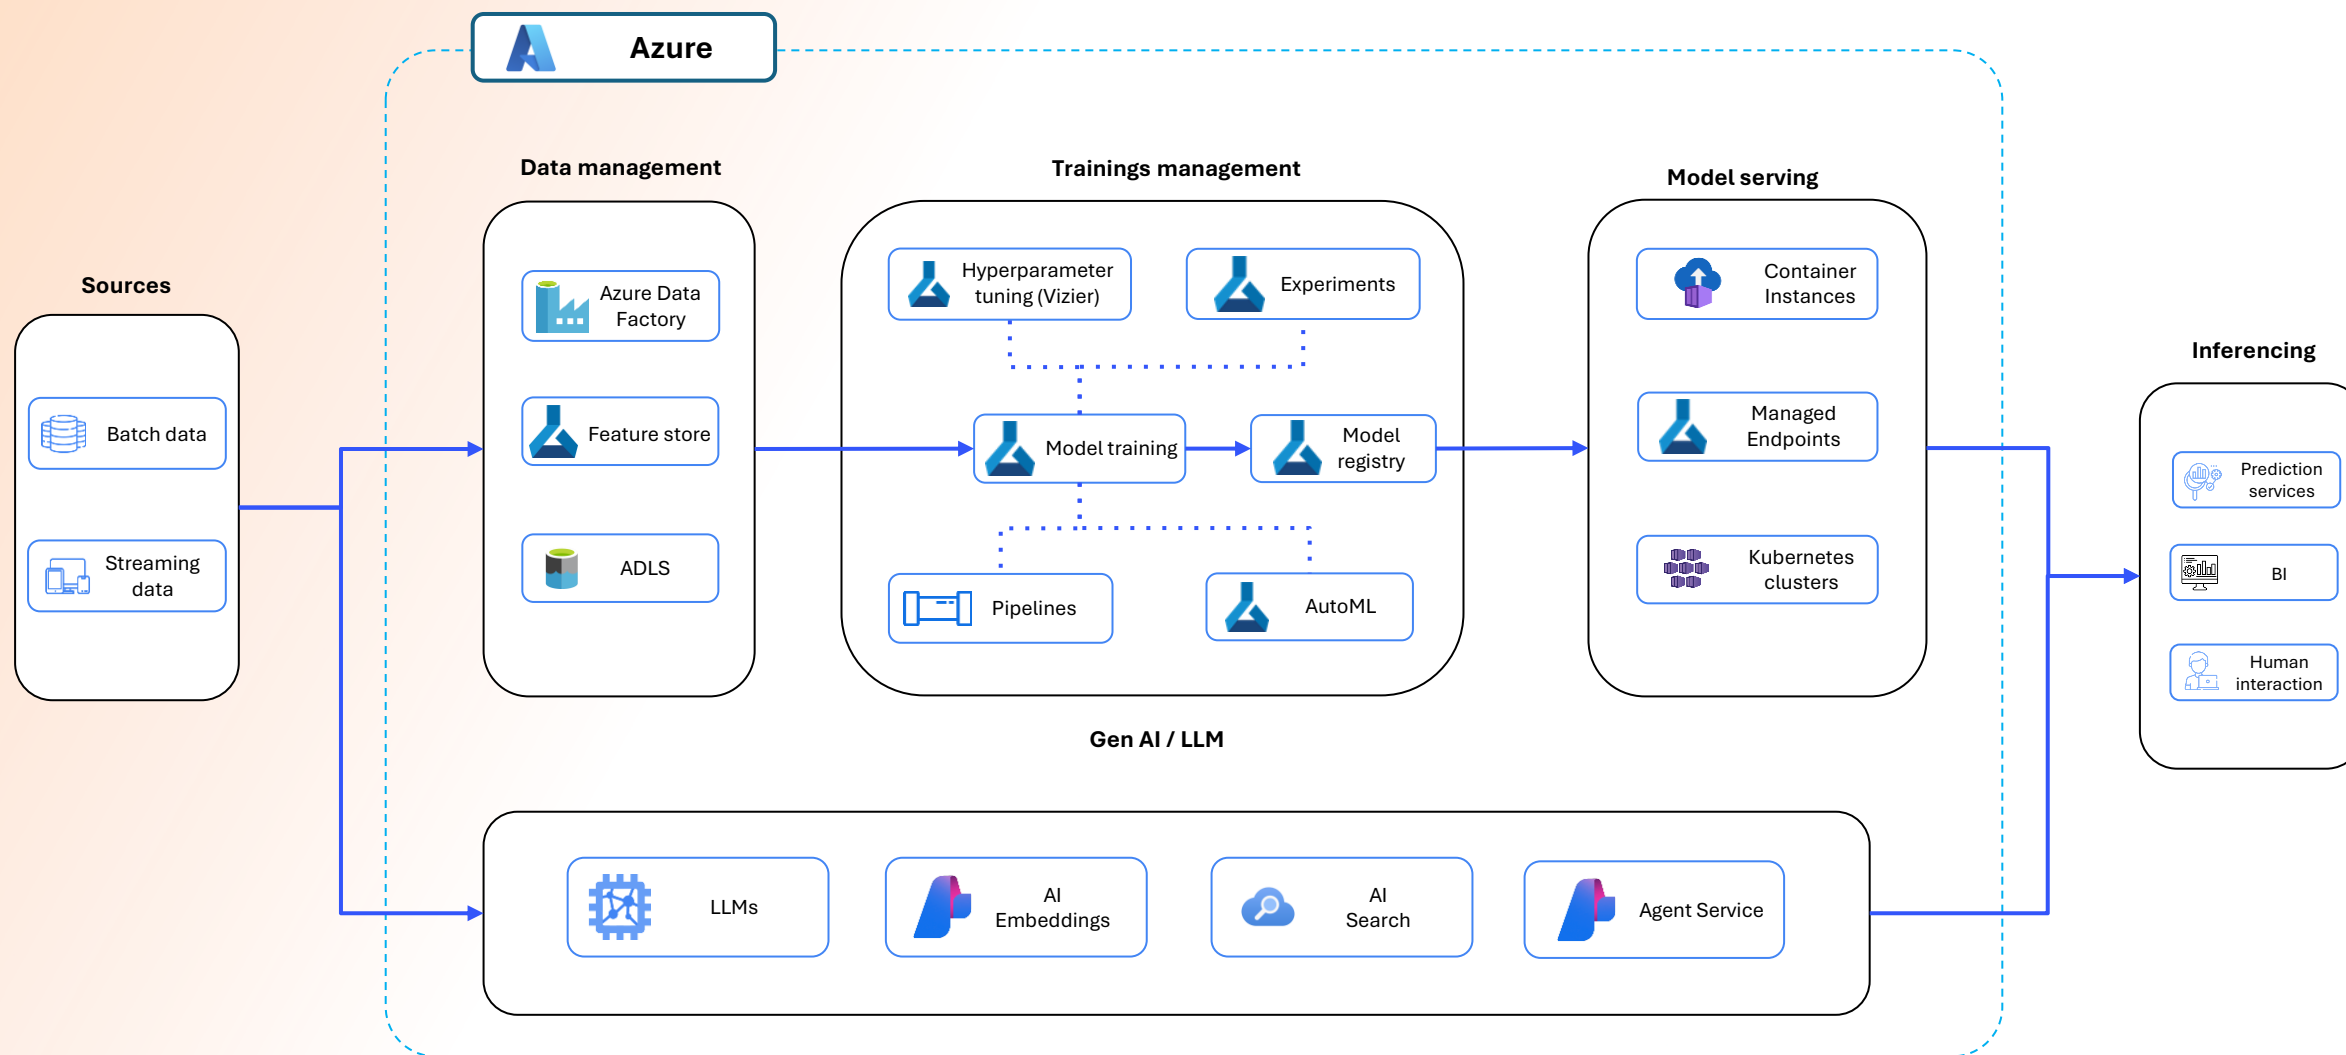
# Division of roles in AI/ML teams

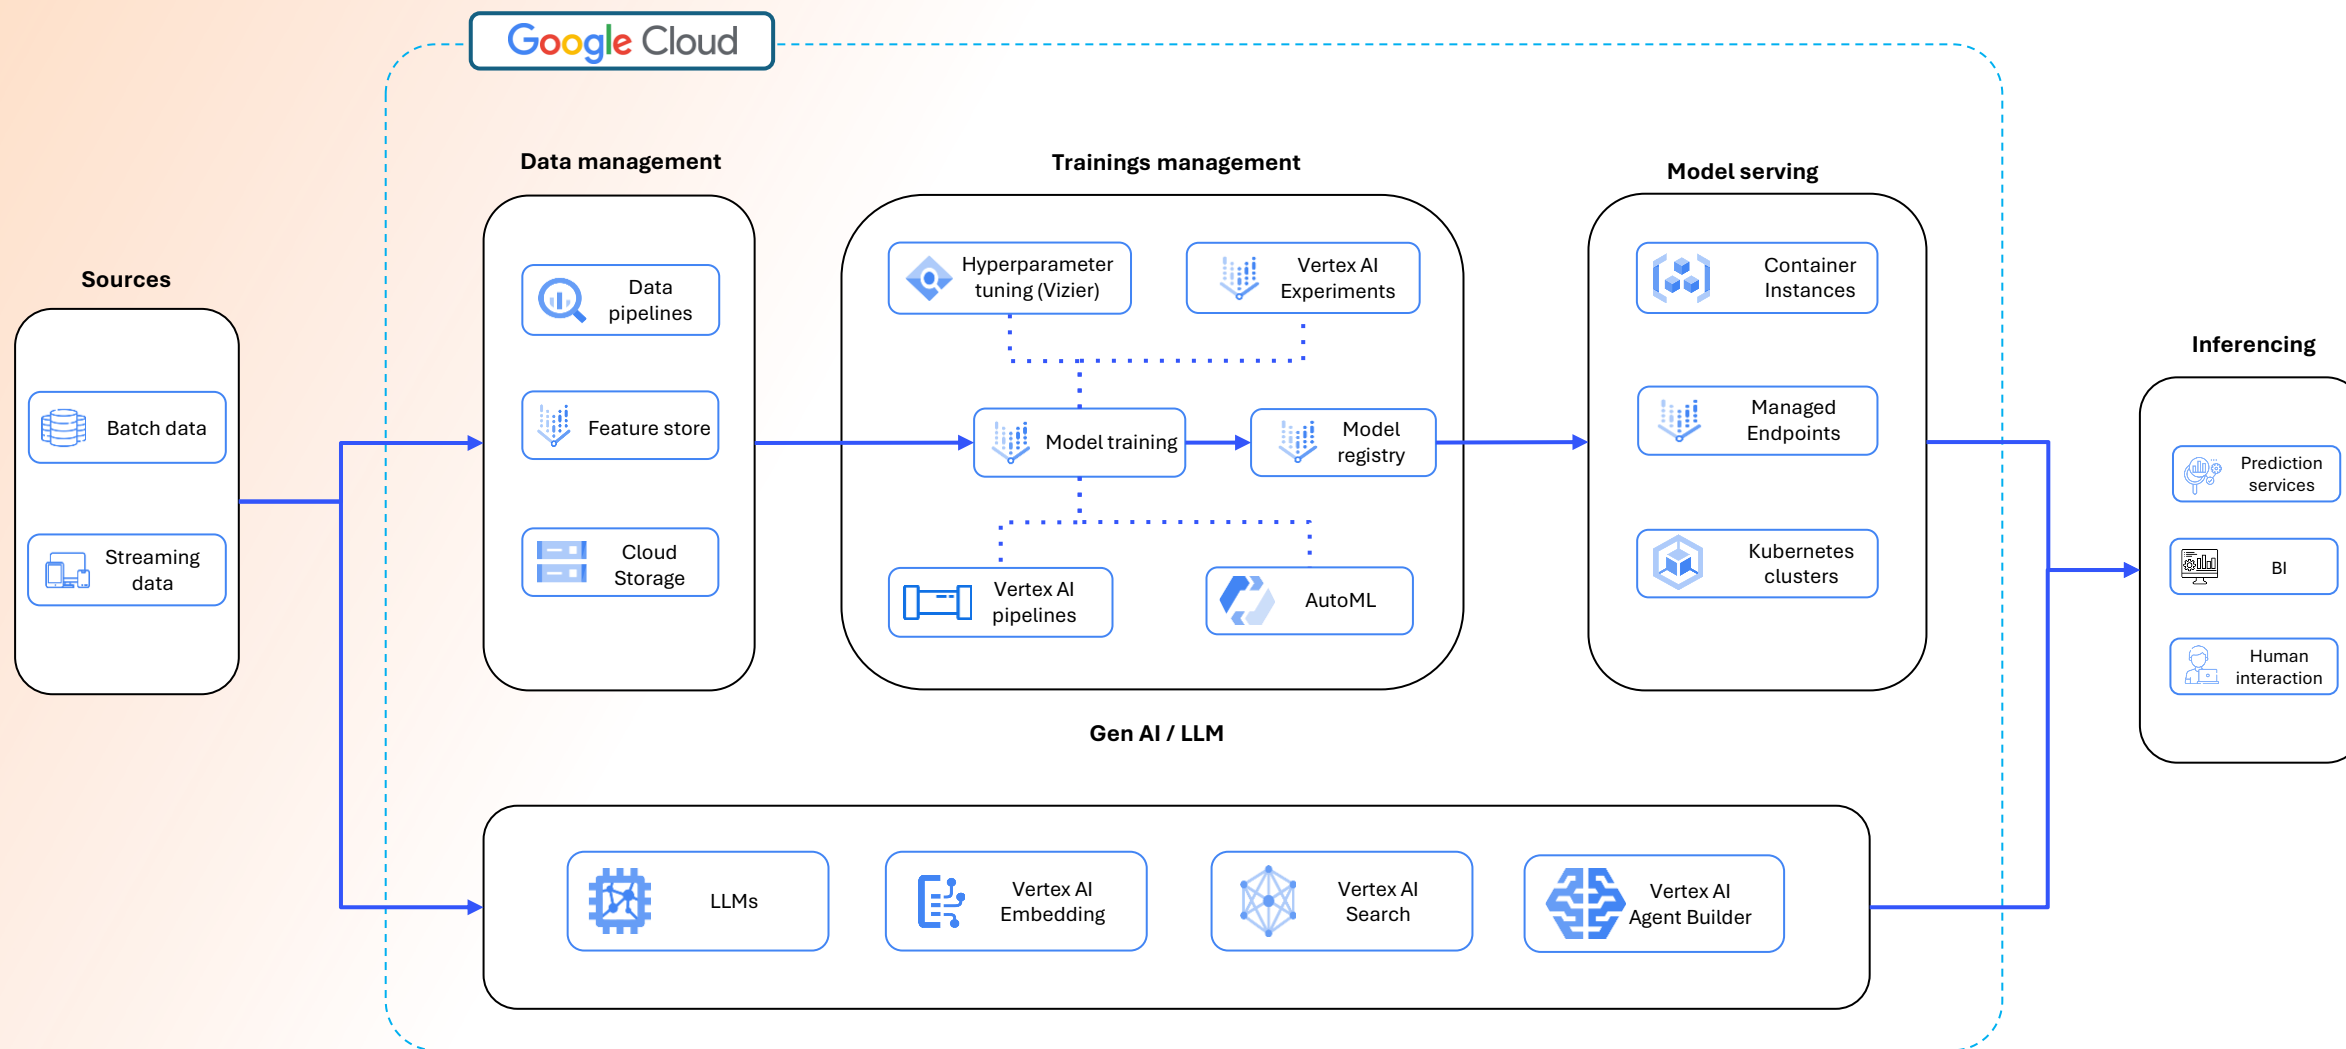


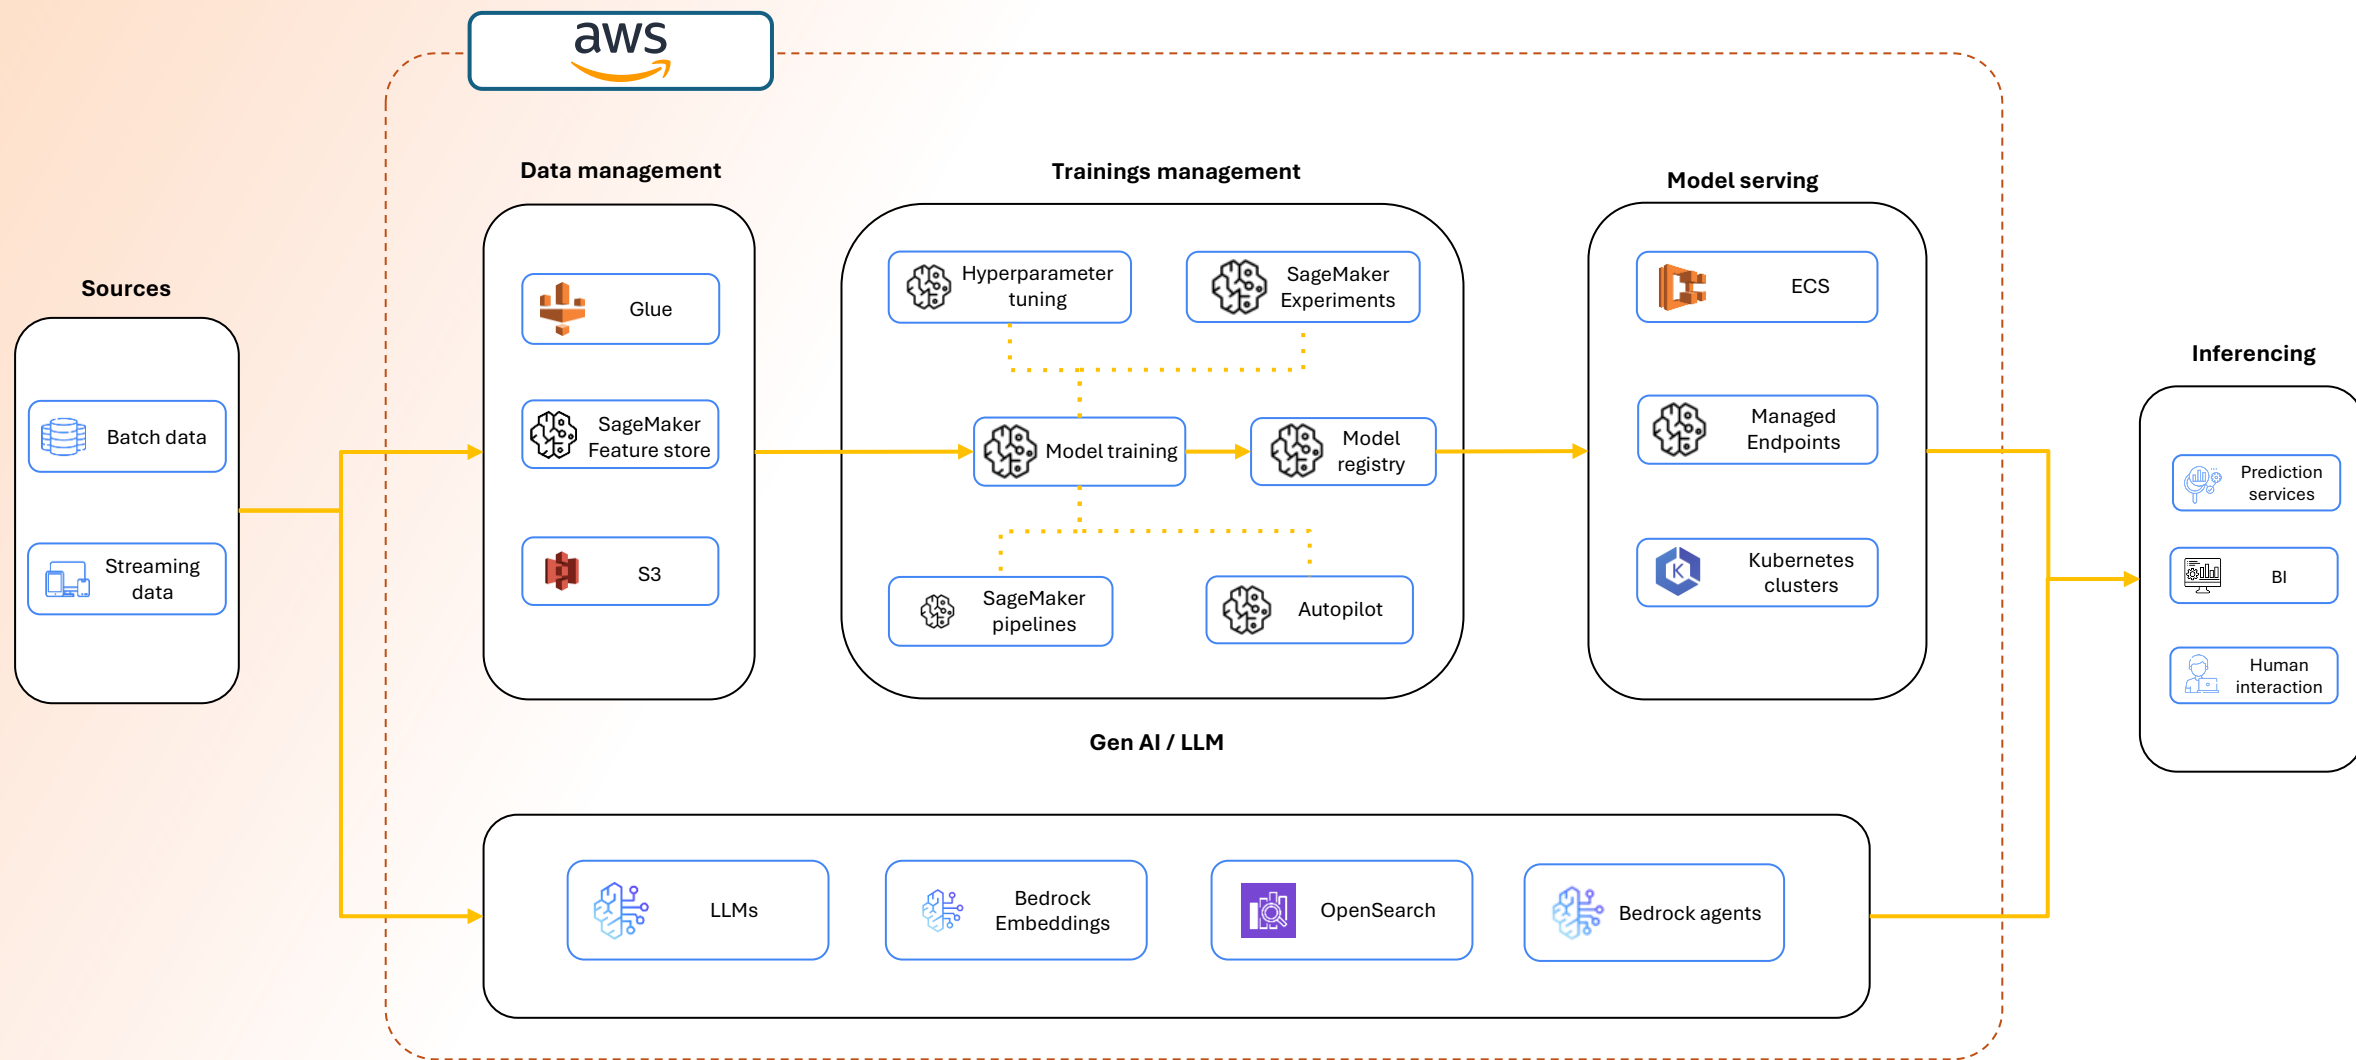


# Cloud Architecture










# DEMO





# System design best practices

## Takeaways

- **Don't overengineer from day one:** start with a simple, working system, add the more advanced features later (if you really need them).
- **If its not tracked, it doesn't exist:** every experiment, model version, and hyperparameter – tracking is never optional.
- **Understand your problem:** AI projects fail due to too high expectations and problem misunderstanding
- **Clean data beats fancy models:** A simple model on good data will outperform any complex architecture trained on garbage – every single time.
- **Choosing your cloud provider matters less than you expect:** the tools differ; the principles don't. Pick one and focus on building value, not migrating clouds.

# Thank you!

**Contact:**



<https://www.linkedin.com/in/maciej-kepa>



<https://github.com/maciejkepa/ai-ml-in-practice>

